

# Basics of Parallel Programs, v.2

## CS 536: Science of Programming, Fall 2022

### A. Why?

- Parallel programs are more flexible than sequential programs but their execution is more complicated.
- Parallel programs are harder to reason about because parts of a parallel program can interfere with other parts.
- Evaluation graphs can be used to show all possible execution paths for a parallel program.

### B. Objectives

After this class, you should know

- The syntax and operational & denotational semantics of parallel programs.

### C. Basic Definitions for Parallel Programs

- **Syntax** for parallel statements:  $S := [S \parallel S \parallel \dots \parallel S]$ . We say  $[S_1 \parallel S_2 \parallel \dots \parallel S_n]$  is the **parallel composition** of the **threads**  $S_1, S_2, \dots, S_n$ .
  - The threads must be sequential: You can't nest parallel programs. (But you can embed parallel programs within otherwise-sequential programs, such as in the body of a loop.)
- **Example 1:**  $[x := x+1 \parallel x := x*2 \parallel y := x^2]$  is a parallel program with three threads. Since it tries to nest parallel programs,  $[x := x+1 \parallel [x := x*2 \parallel y := x^2]]$  is illegal.

### Interleaving Execution of Parallel Programs

- We run sequential threads in parallel by **interleaving** their execution. I.e., we interleave the operational semantics steps for the individual threads.
- We execute one thread for some number of operational steps, then execute another thread, etc.
- Depending on the program and the sequence of interleaving, a program can have more than one final state (or cause an error sometimes but not other times).
- As an example, since evaluation of  $[x := x+1 \parallel x := x*2]$  is done by interleaving the operational semantics steps of the two threads, we can either evaluate  $x := x+1$  and then  $x := x*2$  or evaluate  $x := x*2$  and then  $x := x+1$ .
- The difference between  $[x := x+1 \parallel x := x*2]$  and **if**  $T \rightarrow x := x+1 \square T \rightarrow x := x*2$  **fi** is that the nondeterministic **if-fi** executes only one of the two assignments whereas the parallel composition executes both assignments but in an unpredictable order. The sequential nondeterministic **if-fi**

that simulates the parallel assignments is **if**  $T \rightarrow x := x+1; x := x^*2 \square T \rightarrow x := x^*2; x := x+1$  **fi**. It nondeterministically chooses between the two possible traces of execution for the program.\*

- Because of the nondeterminism, re-executions of a parallel program can use different orders. For example, two executions of **while**  $B$  **do**  $[x := x+1 \parallel x := x^*2]$  **od** can have the same sequence or different sequences of updates to  $x$ .

### Difficult to Predict Parallel Program Behavior

- The main problem with parallel programs is that their properties can be very different from the behaviors of the individual threads.
- **Example 2:**
  - $\models \{x = 5\} x := x+1 \{x = 6\}$  and  $\models \{x = 5\} x := x^*2 \{x = 10\}$
  - But  $\models \{x = 5\} [x := x+1 \parallel x := x^*2] \{x = 11 \vee x = 12\}$
- The problem with reasoning about parallel programs is that different threads can **interfere** with each other: They can change the state in ways that don't maintain the assumptions used by other threads.
- Full interference is tricky, so we're going to work our way up to it. First we'll look at simple, limited parallel programs that don't interact at all (much less interfere).
- But before that, we need to look at the semantics of parallel programs more closely.

### D. Semantics of Parallel Programs

- To execute the sequential composition  $S_1; \dots; S_n$  for one step, we execute  $S_1$  for one step.
- To execute the parallel composition  $[S_1 \parallel \dots \parallel S_n]$  for one step, we take one of the threads and evaluate it for one step.

### Operational and Denotational Semantics of Parallel Programs

- **Definition:** Given  $[S_1 \parallel \dots \parallel S_n]$ , for each  $k = 1, 2, \dots, n$ , if  $\langle S_k, \sigma \rangle \rightarrow \langle T_k, \tau_k \rangle$ , then
 
$$\langle [S_1 \parallel \dots \parallel S_n], \sigma \rangle \rightarrow \langle [S_1 \parallel \dots \parallel S_{k-1} \parallel T_k \parallel S_{k+1} \parallel \dots \parallel S_n], \tau_k \rangle$$
- We write  $E$  for sequential thread that has finished execution, so a parallel program that has finished execution is written  $[E \parallel \dots \parallel E \parallel E]$ . We'll treat  $E$  and  $[E \parallel \dots \parallel E \parallel E]$  as being syntactically equal, i.e.,  $E \equiv [E \parallel \dots \parallel E \parallel E]$ .

#### The $\rightarrow^*$ Notation

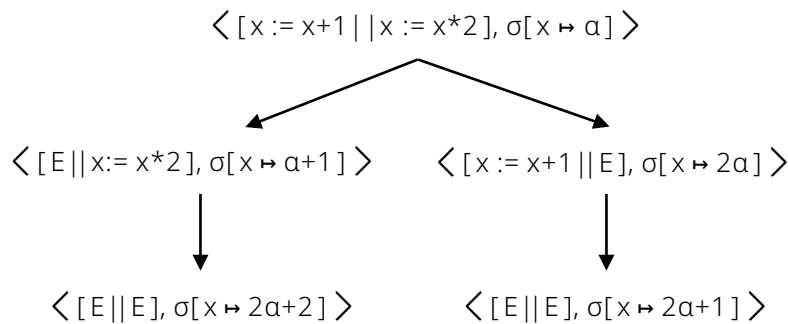
- **Notation:** The  $\rightarrow^*$  notation has the same meaning whether the configurations involved have parallel programs or not:  $\rightarrow^*$  means  $\rightarrow^n$  for some  $n \geq 0$ , where  $C_0 \rightarrow^n C_n$  means that there is actually a sequence of  $n+1$  configurations,  $C_0 \rightarrow C_1 \rightarrow \dots \rightarrow C_{n-1} \rightarrow C_n$  where we've omitted writing the intermediate configurations.
- **Common Mistake:** Writing  $\langle [E \parallel E], \tau \rangle \rightarrow \langle E, \tau \rangle$  is a common mistake. Since  $[E \parallel E] \equiv E$ , going from  $\langle [E \parallel E], \tau \rangle$  to  $\langle E, \tau \rangle$  doesn't involve an execution step. But  $\langle [E \parallel E], \tau \rangle \rightarrow^0 \langle E, \tau \rangle$  is ok.

---

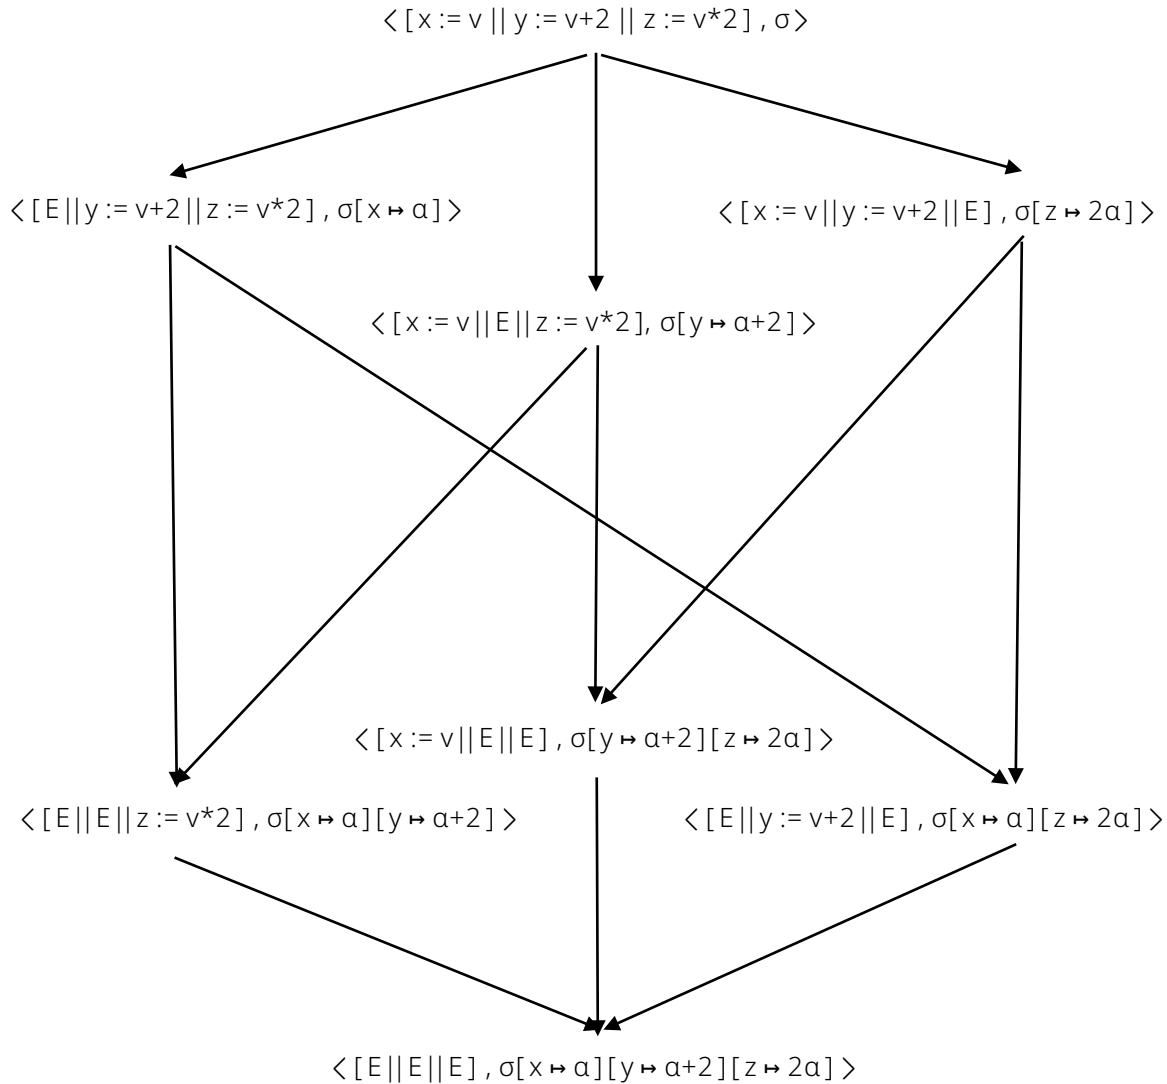
\* This trick doesn't scale up well to larger programs, but it helps with initially understanding parallel execution.

### Evaluation Graph and Denotational Semantics

- Recall that the **evaluation graph** for  $\langle S, \sigma \rangle$  is the directed graph of configurations and evaluation arrows leading from  $\langle S, \sigma \rangle$ .
- When drawing evaluation graphs, the configuration nodes need to be different.
  - (I.e., if the same configuration appears more than once, show multiple arrows into it — don't repeat the same node.)
- An evaluation graph shows all possible executions.
  - A program with  $n$  threads will have  $n$  out-arrows from its configuration.
- A path through the graph corresponds to one possible evaluation of the program.
- The **denotational semantics** of a program in a state is the set of all possible terminating states (plus possibly the pseudostates  $\perp_d$  and  $\perp_e$ ). I.e., the states found in the sinks (i.e., at the leaves) of an evaluation graph. (We'll modify this definition when we get to deadlocked programs.)
  - $M(S, \sigma) = \{\tau \in \sigma \mid \langle S, \sigma \rangle \rightarrow^* \langle E, \tau \rangle\}$ 
    - $\cup \{\perp_d\}$  if  $S$  can diverge; i.e., if  $\langle S, \sigma \rangle \rightarrow^* \langle \perp_d, \tau \rangle$  is possible
    - $\cup \{\perp_e\}$  if  $S$  can produce a runtime error; i.e.,  $\langle S, \sigma \rangle \rightarrow^* \langle \perp_e, \tau \rangle$  is possible
- Example 3:** The evaluation graph below is for the same program as in Example 2, but starting with an arbitrary state  $\sigma$  where  $\sigma(x) = \alpha$ . The graph has two sinks for the two possible final states, so  $M([x := x+1 \parallel x := x*2], \sigma) = \{\sigma[x \mapsto 2\alpha+2], \sigma[x \mapsto 2\alpha+1]\}$ .

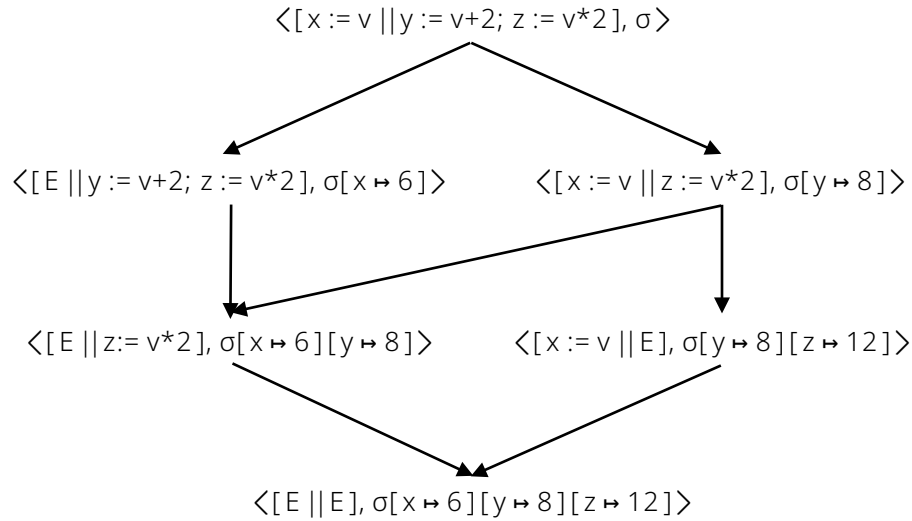


**Example 3**

**Example 4**

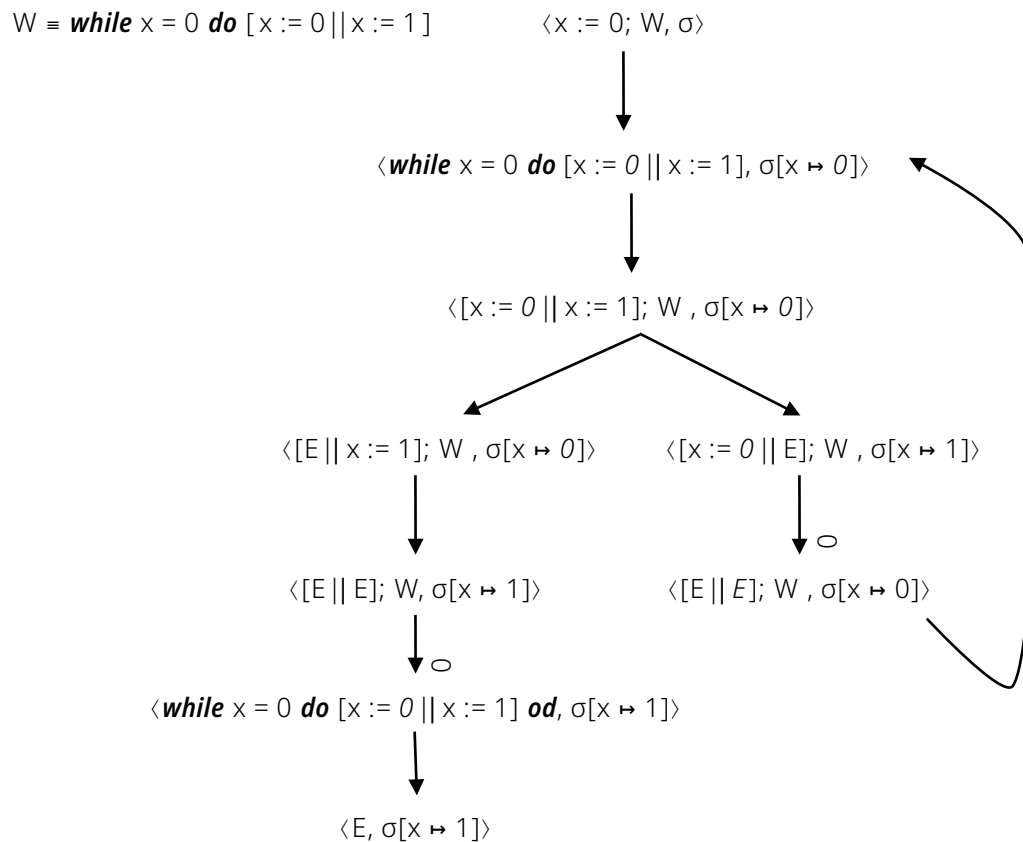
- **Example 4:** For this example, the evaluation graph is for  $\langle [x := v \parallel y := v+2 \parallel z := v*2], \sigma \rangle$ , where  $\sigma(v) = \alpha$ .  $M([x := v \parallel y := v+2 \parallel z := v*2], \sigma) = \{\sigma[x \mapsto \alpha][y \mapsto \alpha+2][z \mapsto 2\alpha]\}$ . Note even though the program is nondeterministic, it produces the same result no matter what execution path it uses. (More generally, if  $S$  is parallel, then  $M(S, \sigma)$  can have more than 1 member, but the converse is not true: Having  $M(S, \sigma)$  of size 1 does not imply that  $S$  is nondeterministic.)

- **Example 5:** If we take the program from Example 4 and combine the last two threads sequentially, then the evaluation graph for the resulting program is a subgraph of the Example 4 graph. Below,  $\sigma(v) = 6$ , and  $M([x := v \parallel y := v+2 \parallel z := v*2], \sigma) = \{\sigma[x \mapsto 6][y \mapsto 8][z \mapsto 12]\}$ .



Example 5

- **Example 6:** Let  $W \equiv x := 0; \text{ while } x = 0 \text{ do } [x := 0 \parallel x := 1] \text{ od}$ . Then  $M(W, \sigma) = \{\sigma[x \mapsto 1], \perp_d\}$ . The problem here is possible divergence, but it only happens if we *always* choose thread 1 when we have to make the nondeterministic choice of  $[x := 0 \parallel x := 1]$ . This is definitely unfair behavior, but it's allowed because of the unpredictability of our nondeterministic choices. In real life, we would want a fairness mechanism to ensure that all threads get to evaluate once in a while.
- If each thread is on a separate processor, then the nondeterministic choice corresponds to which processor is fastest, so the possible divergence of the program is a **race condition**, where the correct behavior of a program depends on the relative speed of the processors involved. Here, divergence occurs if processor 1 is always faster than processor 2 (especially if processor 2 has died).
- Note that it's not necessarily a race condition to have a parallel program producing different results when run multiple times. As long as all results satisfy the specification, there's no race condition.
- **Example 7:** The correctness triple  $\{T\} [x := 0 \parallel x := 1] \{x \geq 0\}$  does not have a race condition, but  $\{T\} [x := 0 \parallel x := 1] \{x > 0\}$  does.

**Example 6**

Note:

- The transitions  $\langle [E \parallel E]; W, \sigma[x \mapsto \dots] \rangle \rightarrow^0 \langle W, \sigma[x \mapsto \dots] \rangle$  take 0 steps because  $[E \parallel E]; W$  and  $W$  are the same program:  $[E \parallel E]; W \equiv E; W \equiv W$ .