

Await and Deadlocks

CS 536: Science of Programming, Fall 2022

A. Why?

- Avoiding interference isn't the same as coordinating desirable activities.
- It's common for one thread to wait for another thread to reach a desired state.
- Care needs to be taken to avoid a program that waits with no hope of completing.

B. Objectives

At the end of this lecture you should know

- The syntax and semantics of the *await* statement.
- How to draw an evaluation diagram for a parallel program that uses *await*.
- How to recognize deadlocked configurations in an evaluation diagram.
- How to list the potential deadlock predicates for a parallel program that uses *await*.

C. Synchronization

The Need for Synchronization

- We've looked at parallel programs whose threads avoid bad interactions.
 - They don't interfere because they don't interact (disjoint programs/conditions).
 - They interact but don't interfere (interference-freedom).
- To supporting good interaction between threads, we often have to have one thread wait for another one. Some examples:
 - Thread 1 should wait until thread 2 is finished executing a certain block of code.
 - Thread 1 has to wait until some buffer is not empty
 - Thread 2 has to wait until some buffer is not full.
- The general problem is that we often want threads to **synchronize**: We want one thread to wait until some other thread makes a condition come true.
- **Example 1**: For a more specific example, in the following program, the calculation of u doesn't start until we finish calculating z , even though u doesn't depend on z .

$$[x := \dots \parallel y := \dots \parallel z := \dots]; u = f(x, y); v := g(u, z)$$

On the other hand, we can't nest parallel programs, so we can't write

$$[[x := \dots \parallel y := \dots]; u = f(x, y) \parallel z := \dots]; v := g(u, z)$$

which would be a natural way to do the calculations of u and z in parallel. In some sense, what we'd like is to run something like

$[x := \dots \parallel y := \dots \parallel \text{wait for } x \text{ and } y; u = f(x, y) \parallel z := \dots]; v := g(u, z)$

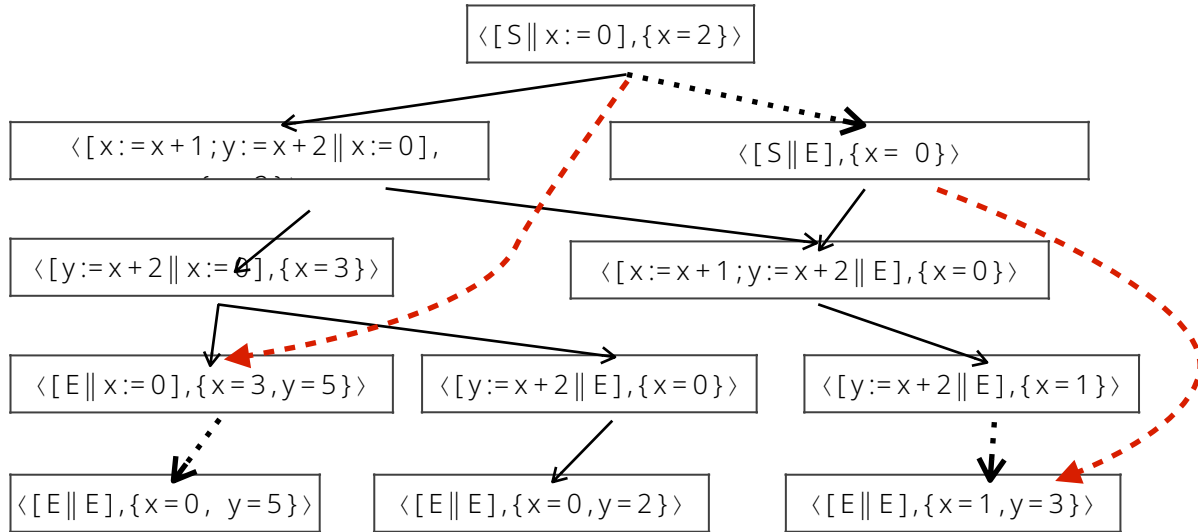
D. The Await Statement

- It's time to introduce a new statement, the **await** statement, whose semantics implements the notion of waiting until some condition is true.
 - Busy wait loops like **while** $\neg B$ **do skip od** $\{B\}$ work but are wasteful.
- **Syntax:** **await** B **then** S **end** where B is a boolean expression and S is a statement.
 - S isn't allowed to have loops, **await** statements, or atomic regions.
 - **await** statements can only appear in sequential threads of parallel programs. (I.e., in some thread S_k in an $[S_1 \parallel S_2 \parallel \dots]$.)
- An **await** statement is a **conditional atomic region**. Suppose that some thread begins with **await** B **then** S **end**, then
 - We nondeterministically choose between all the available threads. I.e., there's no insistence that we must check the **await** before trying other threads. (See case 1 of Example 2.)
 - If we choose the thread that begins with **await** B **then** S **end**,
 - If B is true, then immediately jump to S and execute all of it.
 - The test, jump, and execution of S are atomic — the combination executes as one step. E.g., with the configuration below, we can't set x to 1 between looking up the two x 's to use for calculating $x+x$. (See case 2 of Example 2.)
 - If B is false, we **block**: We wait until B is true. Instead, we nondeterministically choose between the other threads and execute it. (See case 3 of Example 2.)
- An **await** is similar to an atomic **if-then** statement, but not identical.
 - With $\langle \text{if } B \text{ then } S \text{ else skip fi} \rangle$, if B is false, we execute **skip** and complete the **if-fi**. (See case 4, Example 2.)
 - With **await** B **then** S **end**, if B is false, nothing happens until B becomes true. (See the note with case 3, Example 2.)

Example 2:

- (See the discussion above)
 - Case 1: Let $A \equiv \text{await } b \text{ then } x := x+x \text{ end}$ in $\langle [\text{await } b \text{ then } x := x+x \text{ end} \parallel x := 1], \{b = T, x = 2\} \rangle \rightarrow \langle [\text{await } b \text{ then } x := x+x \text{ end} \parallel E], \{b = T, x = 1\} \rangle$.
 - Case 2: $\langle [\text{await } b \text{ then } x := x+x \text{ end} \parallel x := 1], \{b = T, x = 2\} \rangle \rightarrow \langle [E \parallel x := 1], \{b = T, x = 4\} \rangle$. (This is the only transition that executes the **await**.)
 - Case 3: $\langle [\text{await } b \text{ then } S \text{ end} \parallel x := 1], \{b = F\} \rangle \rightarrow \langle [\text{await } b \text{ then } S \text{ end} \parallel E], \{b = F, x = 1\} \rangle$. (The second configuration is blocked, with no other thread available to unblock it.)
 - Case 4: $\langle [\text{if } b \text{ then } x := 0 \text{ else skip fi} \parallel S'], \{b = F\} \rangle \rightarrow \langle [E \parallel S'], \{b = F, x = 0\} \rangle$.

Solid black lines show execution steps taken only when $S \equiv \text{if } x \geq 0 \text{ then } x := x + 1; y := x + 2 \text{ fi}$
Dashed red lines show steps taken only when $S \equiv \text{await } x \geq 0 \text{ then } x := x + 1; y := x + 2 \text{ end}$
Dotted black lines are common to both executions.

Figure 1: Execution of *await* vs *if-fi*

- **Example 3:** Execution of a non-atomic *if-fi* can be interleaved with. In Figure 1, the **dashed red lines** show how execution of *await* $x \geq 0$ *then* $x := x + 1; y := x + 2$ *end* takes just one step to execute the entire body
- **Example 4:** In the introduction, we looked at a situation where we want to wait for some calculations to finish before stating others.

$$[x := \dots \parallel y := \dots \parallel \text{wait for } x \text{ and } y; u = f(x, y) \parallel z := \dots]; v := g(u, z)$$

can be implemented using

```

x_done := F; y_done := F;
[ x := ... ; x_done := T || y := ... ; y_done := T
|| await x_done ∧ y_done then u := f(x, y) end || z := ... ];
v := g(u, z);

```

E. *await*, *wait*, *if*, and $\langle S \rangle$

The Abbreviations $\langle S \rangle$ and *wait B*

- With *await* B *then* S *end*, there are two simple cases: When B is trivial and when S is trivial.

- **Definition:** We can redefine $\langle S \rangle$ to stand for *await* T *then* S *end*. When the test is trivial, we don't need to wait, we simply execute the body atomically. So atomic execution is just conditional atomic execution with a trivial test.
- **Definition:** *wait* $B \equiv$ *await* B *then skip end*. When the body is trivial, we simply wait; when B is true, execution is complete.
- There's an important difference between *wait* B ; S and *await* B *then* S *end*.
 - With *await* B *then* S *end*, once B is true, we immediately atomically execute S , so no other statement can interleave between the test and running S . Therefore S can rely on B being true when it starts executing. If $\sigma(B) = T$, then $\langle [\text{await } B \text{ then } S \text{ end} \parallel \dots], \sigma \rangle \rightarrow \langle [E \parallel \dots], \tau \rangle$, where $\tau \in M(S, \sigma)$.
 - *wait* B ; S means *await* B *then skip end*; S , so it allows another thread to be executed after the *wait* but before running S . If $\sigma(B) = T$, then $\langle [\text{wait } B; S \parallel \dots], \sigma \rangle \rightarrow \langle [S \parallel \dots], \sigma \rangle \rightarrow^* \langle [E \parallel \dots], \tau \rangle$ (if no interleaving occurs). Since interleaving can occur, we rely on B being true when S starts execution.

F. Await Statement Proof Rule and Outlines

- The proof rule for the *await* statement is similar to an *if fi*, but there's no false clause (not even *else skip*).

await Statement (a.k.a. Synchronization Rule)

1. $\{p \wedge B\} S \{q\}$
2. $\{p\} \text{await } B \text{ then } S \text{ end } \{q\}$ *await*, 1

- **Minimal Proof Outline:** $\{p\} \text{await } B \text{ then } S \text{ end } \{q\}$
- **Full Proof Outline:** $\{p\} \text{await } B \text{ then } \{p \wedge B\} S^* \{q\} \text{end } \{q\}$ where S^* is a full proof outline for S .
- **Weakest Preconditions:** $\text{wp}(\text{await } B \text{ then } S \text{ end}, q) \equiv B \rightarrow \text{wp}(S, q)$.
 - This guarantees $\{B \rightarrow \text{wp}(S, q)\} \text{await } B \text{ then } \{ \text{wp}(S, q) \} S^* \{q\} \text{end } \{q\}$
- Note: It may be tempting to write $\{p \wedge \neg B\} \text{await } B \text{ then } \dots$, but that's guaranteed to self-deadlock; the outline is

$$\{p \wedge \neg B\} \text{await } B \text{ then } \{p \wedge \neg B \wedge B\} S^* \{q\} \text{end } \{q\}$$

G. The Producer/Consumer Problem

- The **Producer/Consumer Problem** (a.k.a. **Bounded Buffer Problem**) is a standard problem in parallel programming.
- We have two threads running in parallel: The producer creates things and puts them into a buffer; the consumer removes things from the buffer and does something with them.
- The problem is that if the buffer is full, the producer shouldn't add anything to the buffer; if the buffer is empty, the consumer shouldn't remove anything from the buffer.

- **Example 5:** The rough code to solve this problem is

```

Initialize(buffer);
[while ¬done do                                // Producer
    created := Create();
    await NotFull(buffer) then
        BufferAdd(buffer, created)
    end
od
|| while ¬done do                                // Consumer
    await NotEmpty(buffer) then
        removed := BufferRemove(buffer);
    end;
    Consume(removed)
od
]
```

- Buffer operations need to be synchronized because the threads share the buffer. The threads don't share the created or removed objects, so the Create and Consume calls can go outside the *await* and interleave execution.

H. Deadlock

Blocked Threads; Deadlock

- Recall that $\langle [\text{await } B \text{ then } S \text{ end} ; \dots \parallel \dots], \sigma \rangle$ is blocked (must wait) if $\sigma(B) = F$.
 - If some other thread can make B true, then the *await* may eventually unblock.
 - E.g., $\langle [\text{await } x > y \text{ then } S \text{ end} ; \dots \parallel \dots ; x := y + 1 ; \dots], \sigma \rangle$ could unblock.
 - But if all the other threads have either completed or are themselves blocked, then there's no way for our *await* to unblock. E.g., $\langle [\text{await } B \text{ then } S \text{ end} ; \dots \parallel E], \sigma \rangle$ can't evaluate further.
- **Definition:** A parallel program is **deadlocked** if it has not finished execution and there's no possible evaluation step to take. I.e., all the threads are either complete or blocked and at least one thread is blocked.
- If all the other threads are complete or are blocked, the program is **deadlocked**: There's no possible evaluation step leaving from the configuration.
- **Example 6:** If $A \equiv \text{await } x \geq 0 \dots \text{end}$, then there's no arrow out of $\langle [A \parallel A], \sigma[x \mapsto -1] \rangle$, so this configuration is deadlocked. If the value of x had been ≥ 0 , then both *await* statements would have been eligible for execution. Since only one blocked thread is required for deadlock, $\langle [\text{await } x \geq 0 \dots \text{end} \parallel E], \sigma[x \mapsto -1] \rangle$ is also deadlocked.
- Threads can block themselves (trivial example: *await false then S end*).

- More often, threads block because they're waiting for conditions they expect other threads to establish. E.g., if we're running in a state where $y = 0$ and $x = 0$, then these two threads deadlock:
 - Thread 1: $\{p_1\}$ **await** $y \neq 0$ **then** $x := 1$;...
 - Thread 2: $\{p_2\}$ **await** $x \neq 0$ **then** $y := 1$;...
- A program might deadlock under all execution paths or only certain execution paths.
- **Example 7:** The program

$[\text{await } y \neq 0 \text{ then } x := 1 \text{ end} \parallel \text{await } x \neq 0 \text{ then } y := 1 \text{ end}]$

deadlocks iff you execute in a state where x and y are both zero

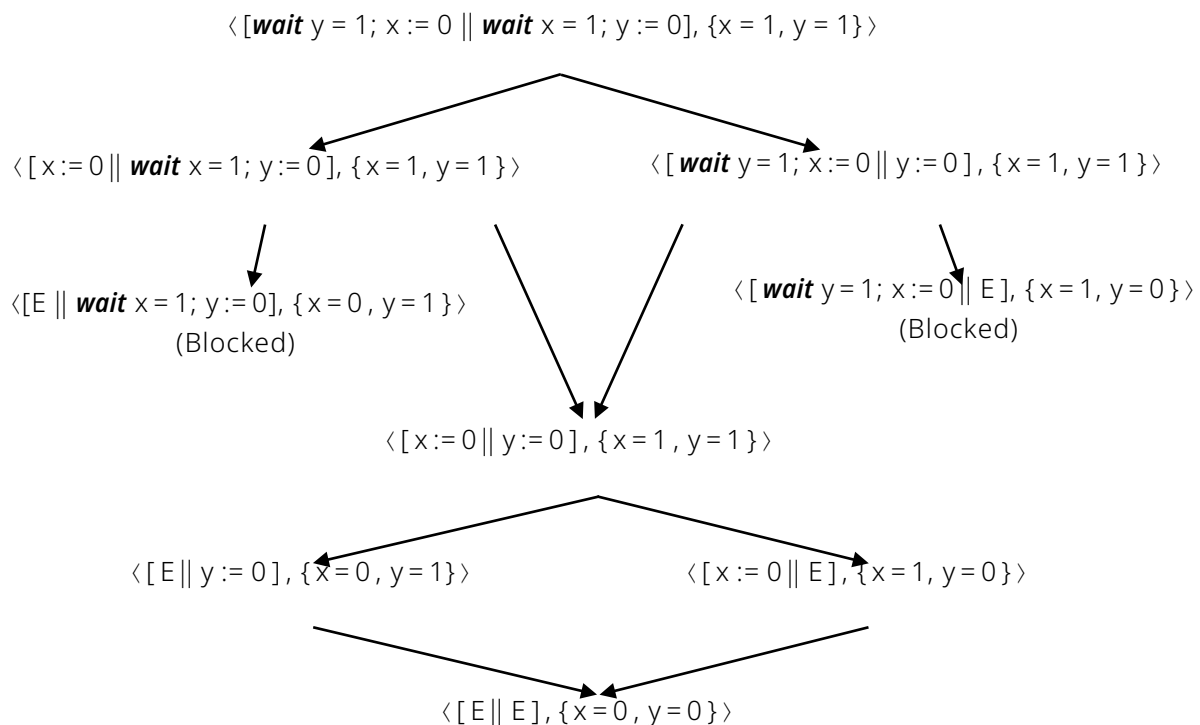


Figure 2: A Program That Can Deadlock

- **Example 8:** If thread 1 sets $x := 0$ before thread 2 evaluates its **wait** x , then thread 2 will block. (Recall **wait** $x \equiv \text{await } x \text{ then skip end}$.)

$\{T\} x := 1; y := 1;$
 $[\text{wait } y = 1; x := 0 \parallel \text{wait } x = 1; y := 0]$
 $\{x = 0 \wedge y = 0\}$

Figure 2 contains an execution graph for this program in state $\{x = 1, y = 1\}$ (somewhat abbreviated). There are two deadlocking paths (and four paths that terminate correctly).

- Obviously, we'd like to know if a program is going to deadlock. The following test identifies a set of predicates that indicate potential problems with a program; if none of these predicates is satisfiable, then deadlock is guaranteed not to occur.
- If one or more of these predicates is satisfiable, then we can't guarantee that deadlock will not occur, but we aren't guaranteeing that deadlock **must** occur. (So the deadlock conditions are sufficient to show deadlock is impossible but they are not necessary conditions.)
- Let $\{p\} [\{p_1\} S_1^* \{q_1\} \parallel \{p_2\} S_2^* \{q_2\} \parallel \dots \parallel \{p_n\} S_n^* \{q_n\}] \{q\}$ be a full outline for a parallel program, where $p \equiv p_1 \wedge \dots \wedge p_n$ and $q \equiv q_1 \wedge \dots \wedge q_n$.
- Definition:** A (**potential**) **deadlock condition** for the program outline above is a predicate of the form $r_1' \wedge r_2' \wedge \dots \wedge r_n'$ where each r_k' is either
 - q_k , the postcondition for thread S_k or
 - $p \wedge \neg B$ where $\{p\}$ **await** B ... appears in the proof outline for thread S_k .
 - In addition, at least one of the r_k' must involve waiting. I.e., $q \equiv q_1 \wedge \dots \wedge q_n$ is not a potential deadlock condition.
- A program outline is **deadlock-free** if every one of its potential deadlock conditions is unsatisfiable (i.e., a contradiction):
 - I.e., for each deadlock condition r' , we have $\models \neg r'$ (or the equivalent $\models r' \rightarrow F$).

Parallelism with Deadlock Freedom

1. $\{p_1\} S_1^* \{q_1\}$
2. $\{p_2\} S_2^* \{q_2\}$
- ...
- n. $\{p_n\} S_n^* \{q_n\}$
- n+1. $\{p_1 \wedge p_2 \wedge \dots \wedge p_n\}$
 $[S_1 \parallel S_2 \parallel \dots \parallel S_n]$
 $\{q_1 \wedge q_2 \wedge \dots \wedge q_n\}$

D.P. w/o deadlock, 1, 2, ..., n

where the $\{p_k\} S_k^* \{q_k\}$ are pairwise interference-free standard proof outlines and the parallel program outline is deadlock-free.

I. Examples of Deadlock Conditions

- Example 9:** Let's take the program from Example 7:

$[\text{await } y \neq 0 \text{ then } x := 1 \text{ end } \parallel \text{await } x \neq 0 \text{ then } y := 1 \text{ end}]$

and develop an annotation for it:

$\{T\}$
 $[\{T\} \text{await } y \neq 0 \text{ then } \{y \neq 0\} x := 1 \{x \neq 0 \wedge y \neq 0\} \text{end } \{x \neq 0 \wedge y \neq 0\}$
 $\parallel \{T\} \text{await } x \neq 0 \text{ then } \{x \neq 0\} y := 1 \{x \neq 0 \wedge y \neq 0\} \text{end } \{x \neq 0 \wedge y \neq 0\}$
 $] \{x \neq 0 \wedge y \neq 0\}$

- Let set $D_1 = \{x \neq 0 \wedge y \neq 0, y = 0\}$ be the choices for p_1' .

- $x \neq 0 \wedge y \neq 0$ is the thread postcondition
- $y = 0$ indicates thread 1 is blocked at the **await** statement.
- Similarly, let set $D_2 = \{x \neq 0 \wedge y \neq 0, x = 0\}$ be the choices for p_2' (the postcondition of thread 2 and the blocking condition for its **await**).
- There are three choices for the potential deadlock predicate $r_1' \wedge r_2'$:
 - $(x \neq 0 \wedge y \neq 0) \wedge (x = 0)$, which is a contradiction.
 - $(y = 0) \wedge (x \neq 0 \wedge y \neq 0)$, which is a contradiction.
 - $(y = 0) \wedge (x = 0)$, which is not a contradiction, therefore, it's a potential deadlock condition, and our program does not pass the deadlock-freedom test.
- Recall $(x \neq 0 \wedge y \neq 0) \wedge (x \neq 0 \wedge y \neq 0)$ is not a potential deadlock predicate because it says that the two threads have both completed.
- One way out of this predicament is to make the initial precondition the negation of $y = 0 \wedge x = 0$. Let p be $(x \neq 0 \vee y \neq 0)$ in

```

{p}
[ {p} await  $y \neq 0$  then { $p \wedge y \neq 0$ }  $x := 1$  { $x \neq 0 \wedge y \neq 0$ } end { $x \neq 0 \wedge y \neq 0$ }
|| {p} await  $x \neq 0$  then { $p \wedge x \neq 0$ }  $y := 1$  { $x \neq 0 \wedge y \neq 0$ } end { $x \neq 0 \wedge y \neq 0$ }
] { $x \neq 0 \wedge y \neq 0$ }

```

- Let $D_1 = \{x \neq 0 \wedge y \neq 0, p \wedge y = 0\}$ and let $D_2 = \{x \neq 0 \wedge y \neq 0, p \wedge x = 0\}$.
- The three potential deadlock predicates are now contradictory
 - $(x \neq 0 \wedge y \neq 0) \wedge (p \wedge x = 0)$ (is false because of $x \neq 0 \wedge x = 0$)
 - $(p \wedge y = 0) \wedge (x \neq 0 \wedge y \neq 0)$ (is false because of $y = 0 \wedge y \neq 0$)
 - $(p \wedge y = 0) \wedge (p \wedge x = 0)$

$$\equiv ((x \neq 0 \vee y \neq 0) \wedge y = 0) \wedge ((x \neq 0 \vee y \neq 0) \wedge x = 0)$$

$$\Rightarrow (x \neq 0 \wedge y = 0) \wedge (y \neq 0 \wedge x = 0)$$

$$\Rightarrow F$$
- (end of example 9)
- **Example 10:** Since it has three threads, the deadlock conditions for this program are a bit more involved than for Example 9. Thread 1 has one **await** statement, thread 2 has two **await** statements, and thread 3 has no **await** statements.

```

[ ... { $p_{11}$ } await  $B_{11}$  ... { $q_1$ }
|| ... { $p_{21}$ } await  $B_{21}$  ... { $p_{22}$ } await  $B_{22}$  ... { $q_2$ }
|| ... { $q_3$ } ]

```

- The deadlock conditions are built using the three sets
 - $D_1 = \{p_{11} \wedge \neg B_{11}, q_1\}$
 - $D_2 = \{p_{21} \wedge \neg B_{21}, p_{22} \wedge \neg B_{22}, q_2\}$
 - $D_3 = \{q_3\}$.

- Let D be the set of deadlock conditions, $D = \{r_1 \wedge r_2 \wedge r_3 \mid r_1 \in D_1, r_2 \in D_2, r_3 \in D_3\} - \{q_1 \wedge q_2 \wedge q_3\}$. Specifically, we get the following ($2 \times 3 \times 1 - 1 = 5$) conditions:

$D = \{ (p_{11} \wedge \neg B_{11}) \wedge (p_{21} \wedge \neg B_{21}) \wedge q_3,$	— Thread 1 blocked; thread 2 blocked at 1st await
$(p_{11} \wedge \neg B_{11}) \wedge (p_{22} \wedge \neg B_{22}) \wedge q_3,$	— Thread 1 blocked; thread 2 blocked at 2nd await
$(p_{11} \wedge \neg B_{11}) \wedge q_2 \wedge q_3,$	— Thread 1 blocked
$q_1 \wedge (p_2 \wedge \neg B_{21}) \wedge q_3,$	— Thread 2 blocked at 1st await
$q_1 \wedge (p_{22} \wedge \neg B_{22}) \wedge q_3 \}$	— Thread 2 blocked at 2nd await
- The program will be deadlock-free if every predicate in D is a contradiction (i.e., unsatisfiable).

J. Strengthening Deadlock Conditions

- Having all deadlock conditions be contradictory is sufficient for guaranteeing that no program execution will deadlock.
- It's not a necessary condition, however. Just because some $r \in D$ is satisfiable, that doesn't mean that there exists a program execution that can get to the corresponding deadlocked configuration.

- Example 11:** Here's an example of strengthening conditions so that we can prove deadlock freedom. The program is small enough for us to be able to hand-verify that it never deadlocks (by figuring out all possible interleavings).

$\{T\} n := 0; [\text{await } n = 0 \text{ then } n := 1 \text{ end} \parallel \text{wait } n = 1] \{n > 0\}$

- If we annotate the program as below, we have sequential correctness for each thread, plus the threads are interference-free:

$\{T\} n := 0; \{T\}$
 $[\{T\} \text{ await } n = 0 \text{ then } n := 1 \text{ end } \{n > 0\}$
 $\parallel \{T\} \text{ wait } n = 1 \{n > 0\}$
 $]\{n > 0\}$

- On the other hand, we can't prove deadlock freedom. There are $2 \times 2 - 1 = 3$ deadlock conditions and all of them are satisfiable:

- $n \neq 0 \wedge n \neq 1$ — Both threads blocked
- $n \neq 0 \wedge n > 0$ — Thread 1 blocked
- $n > 0 \wedge n \neq 1$ — Thread 2 blocked

- The problem here is that the proof outline's conditions are too weak. We want each deadlock condition to be logically equivalent to false, the strongest predicate.
- To make a conjunctive formula stronger, we need to strengthen its conjuncts. For a deadlock-freedom test, we have two kinds of conjuncts:
 - (postcondition of thread)
 - (precondition of *await* statement) $\wedge \neg$ (test of *await* statement)

- By strengthening the postcondition of the initial assignment of $n := 0$ from true to $n = 0$, we can strengthen the precondition of the first **await**:

```
{T} n := 0; {n = 0}
[ {n = 0} await n = 0 then {n = 0 ∧ n = 0} n := 1 end {n > 0}
|| {T} await n = 1 then {n = 1} skip {n = 1} end {n > 0} ]
{n > 0}
```

- The potential deadlock conditions for the proof outline above are now
 - $(n = 0 \wedge n \neq 0) \wedge n \neq 1$ — Both threads blocked (contradiction)
 - $(n = 0 \wedge n \neq 0) \wedge n > 0$ — Thread 1 blocked (contradiction)
 - $n > 0 \wedge n \neq 1$ — Thread 2 blocked (satisfiable)
- So two of the conditions are contradictory, but one condition is still satisfiable. To prove deadlock-freedom, we need to strengthen the conditions even more to include the state we get to when the first thread has executed and the second thread hasn't.

- Unfortunately, if we annotate the two threads as

- $\{n = 0\} \text{ **await** } n = 0 \text{ **then** } n := 1 \text{ **end** } \{n = 1\}$
- $\{n = 1\} \text{ **wait** } n = 1 \{n = 1\}$

- Then the precondition of the parallel program has to be $(n = 0) \wedge (n = 1)$, which doesn't follow from the strongest postcondition of $n := 0$ and worse yet, isn't possible anyway.

```
{T} n := 0;
{n = 0 ∧ n = 1} // ← error
[ {n = 0} await n = 0 then n := 1 end {n = 1}
|| {n = 1} wait n = 1 {n = 1}
] {n = 1 ∧ n = 1} {n = 1}
```

- Before thread 2 runs, it sees $n = 0$ or $n = 1$ depending on whether thread 1 has run yet. If we use that as the precondition for thread 2, then we get $n = 0 \wedge (n = 0 \vee n = 1)$ as the precondition for the parallel program, which works:

```
{T} n := 0;
n = 0 ∧ (n = 0 ∨ n = 1)
[ {n = 0} await n = 0 then n := 1 end {n = 1}
|| {n = 0 ∨ n = 1} wait n = 1 {n = 1}
] {n = 1 ∧ n = 1} {n = 1}
```

- Better still, the deadlock conditions are now all contradictions, so we have deadlock-freedom

- $(n = 0 \wedge n \neq 0) \wedge ((n = 0 \vee n = 1) \wedge n \neq 1)$ — Both blocked (contradiction)
- $(n = 0 \wedge n \neq 0) \wedge n = 1$ — Thread 1 blocked (contradiction)
- $n = 1 \wedge ((n = 0 \vee n = 1) \wedge n \neq 1)$ — Thread 2 blocked (contradiction)

- Unfortunately, one of the interference freedom tests now fails:

- Pass: $\{n = 0 \wedge (n = 0 \vee n = 1)\} \text{ **await** } n = 0 \text{ **then** } n := 1 \text{ **end** } \{n = 0 \vee n = 1\}$
- Pass: $\{n = 0 \wedge n = 1\} \text{ **await** } n = 0 \text{ **then** } n := 1 \text{ **end** } \{n = 1\}$

- Fail: $\{(n=0 \wedge n=1) \wedge n=0\} \text{ *wait* } n=1 \{n=0\}$
- We can solve this problem by adding an auxiliary variable to say whether or not the first thread has run and set $n=1$. (end of Example 11)